

The market is a column nobody reads

Market & geo differentiation, the language layer, prompt architecture and the custom-agent surface — audited against the code, the production database, and primary market sources.

DATE	10 August 2026
REPO / COMMIT	Aurora-091/openvent · main @ 8498e58
DATA	Production Postgres (Supabase, pooled) — 4 orgs, 17 agent configs, 9 templates, 25 calls
CONTENTS	Audit 15 (9 findings) · ADR-095 (proposed)
STATUS	Findings only — no code changed, nothing committed

Read the citation labels. Every WER figure is cited to arXiv 2604.19151v2 and was *not* measured on Weeber audio; every price is vendor- or community-published. The TRAI Direction and the arXiv paper were fetched directly. Production behaviour is inferred from `main` plus DB rows — the deployed commit cannot be fingerprinted (no health route). See §8.

Audit 15 — the market is a column nobody reads

Audits 13 and 14 asked whether a call is fast and whether the agent knows what it is saying. This one asks the question one level up: **the code says "India-first", the deployment is entirely US, and the two are wired together by nothing.** `agent-frame.ts:53` states it outright in a comment — "India-first list since that's the primary market" — and `pricing-lock-2026-07-18.md` decision #3 locked "India and Global are separate plans, not a currency toggle." Neither statement has a field, a code path, or a production row behind it.

Everything below is either read out of the source at `8498e58` , queried out of the production database, or cited to a primary external source with the citation attached. Where a number is vendor-published or paper-published rather than measured on Weeber's own audio, it says so. Nothing here was benchmarked locally — there are no provider keys in this environment (see §8).

0. The headline

Four findings decide the next quarter, and three of them are the same defect wearing different clothes: **market is treated as a property of the phone number being dialled, not a property of the business.**

#	FINDING	SEV
F1	The India insurance vertical is legally undeliverable as built, and our own gate correctly blocks it. TRAI's 16 Dec 2025 Direction bars IRDAI-regulated entities from <i>any</i> service/transactional call on a non-1600-series number after 15 Feb 2026 — a deadline that has passed. All four production numbers are US Twilio DID's with <code>number_series</code> NULL.	P0
F2	<code>orgs.country_code</code> / <code>currency</code> / <code>timezone</code> are write-only . They are accepted by <code>PATCH /api/app/settings</code> , echoed by <code>GET /me</code> , and read by zero code in <code>packages/api/src/voice/</code> . Jurisdiction is inferred at dial time from the <i>recipient's</i> phone prefix instead. 3 of 4 production orgs have all three NULL.	P0
F3	The entire Indic investment is unreached in production. <code>language</code> is NULL on 14 of 17 <code>org_agent_configs</code> ; <code>buildLanguageInstructionBlock</code> returns <code>" "</code> for null/ <code>en</code> ; Sarvam ran on 0 of 25 production calls. The ₹1.2/min Sarvam COGS that the locked India margin's <i>good</i> case depends on has never occurred.	P1
F4	The production STT default (Deepgram Nova-3) is rated Tier III — "inadequate" — on Indian telephonic speech by the largest published benchmark of it, with most Indic languages unsupported outright. The Tier I model on that benchmark (Sarvam) is fully wired, program-credited, and idle.	P1
F5	<code>language</code> is a free-text input on a field that silently switches STT provider, TTS provider, voice identity, and disclosure language. <code>"hi-IN"</code> and <code>"Hindi"</code> both fail <code>prefersSarvam()</code> and fall back to the English-first stack with no warning.	P1
F6	Localised recording disclosures exist for 3 of 12 recommended languages (<code>en</code> , <code>hi</code> , <code>hinglish</code>). A Tamil agent speaks a Sarvam Tamil voice and an English legal disclosure. The <code>hi / hinglish</code> lines are marked in-code as drafts pending human review — and are the live path.	P1
F7	"Custom agents" is an admin-only authoring surface. 9 templates, all <code>visibility=public</code> , <code>owner_org_id</code> NULL, zero private. <code>persona_prompt</code> is NULL/empty on 17 of 17 configs — no customer has ever customised a prompt. That is a services motion priced as self-serve SaaS.	P1
F8	<code>calls</code> has no <code>template_key</code> , no <code>agent_config_id</code> , no <code>language</code> . A call cannot be attributed to an agent, a template, or a language — so F3, F4, F6 and F7 are unanswerable from data, permanently, by design.	P1

#	FINDING	SEV
F9	The US-only, English-only, 19,480-character final-expense template is enabled on 2 of 3 insurance orgs. Market scope for a template lives only in prose inside <code>default_persona_prompt</code> ; no <code>region / market</code> column exists on <code>agent_templates</code> or <code>org_agent_configs</code> .	P2

1. F1 (P0) — India insurance cannot legally be delivered, and we block it ourselves

The rule (primary source, fetched)

TRAI Direction dated 16 December 2025, F. No. G-6/(8)/2025-QoS-Part(I) (E-18071), retrieved from https://www.trai.gov.in/sites/default/files/2025-12/Direction_16122025.pdf. Operative content:

- 1600-series adoption by all IRDAI-regulated entities was to be **completed by 15 February 2026**.
- Failure attracts action "as per regulatory provisions applicable to **unregistered telemarketers**" under TCCCPR 2018 reg. 3 — i.e. suspension/disconnection of telecom resources.
- Clause (iii), verbatim: *"entities regulated by IRDAI are not permitted to initiate any service or transactional voice calls, even with the explicit or inferred consent of customers, from numbers other than those allocated under the 1600-series, after the 15th February 2026."*
- 15-day status reporting obligation.

Note the scope of clause (iii): **consent does not cure it**. This is not a promotional-calling rule that a consent record satisfies. It is a number-origination rule. Our consent machinery — which is genuinely good — is irrelevant to it.

The repo already knew this. `docs/agent-prompts/00-insurance-regulatory-reference.md:21` describes the mandate accurately and flags it as a real platform gap; the gap was subsequently built.

What the code does

`packages/api/src/voice/compliance/insurance-gates.ts:74` requires an **active `org_phone_numbers` row with `number_series = "1600"`** for any insurance org dialling India. It is unconditional — no feature flag, no vertical exemption beyond `orgs.vertical === "insurance"`. It is wired into both dial paths: `workflows/scheduler.ts:81` (`reason: "insurance_number_series"`) and `voice/routes.ts:289` (403).

The general non-insurance counterpart, `compliance/number-series-gate.ts`, requires a 140/160-series number for other verticals — but it is behind `INDIA_NUMBER_SERIES_FLAG`, **off by default**, with the reason documented in the file header: a live Shopify org dials India numbers today from an unregistered Twilio number, and turning the gate on unconditionally would have broken it.

What production actually holds

id	org_id	phone_number	number_series	status	provider
4	org_58c7d5cc-... (krisn, insurance)	+17126257861	(null)	active	twilio
5	org_68497dd7-... (rishipawar8999, ins.)	+17754554413	(null)	active	twilio
6	org_58c7d5cc-...	+17744745274	(null)	released	twilio
7	org_a4ddb581-... (presistentads, ins.)	+18573706834	(null)	active	twilio

Four numbers. All `+1`. All Twilio. `number_series` NULL on every row. Three of the four production orgs are `vertical = insurance`, and one of them (`org_68497dd7`) is the only org with `country_code = IN` / `timezone = Asia/Kolkata` set

at all.

The conclusion

For every insurance org in production, an outbound call to any `+91` number is blocked at dial time and always will be, because:

1. The gate demands a 1600-series number.
2. Twilio does not allocate 1600-series numbers. That series is allocated to IRDAI-regulated entities through Indian TSPs and coordinated by Indian CPaaS vendors — Exotel, Ozonetel, FreJun, C-Zentrix (vendor-published onboarding claims; Exotel advertises 1600-series plus DLT registration in under 48 hours — vendor claim, not verified by us).
3. `project-brief.md` lists Exotel as supported BYO India telephony. **No production org uses it.** Every row above says `twilio`.

The gate is correct and I would not weaken it. The finding is not "the gate is wrong" — it is that **the gate is the only part of the India insurance story that exists**. The compliance check for a capability we do not have is a very expensive way to discover we do not have it.

Unverified, do not repeat: a vendor blog claims these rules carry ₹10 lakh penalties. No primary source supports that. The verified sanction is unregistered-telemarketer treatment under TCCCPR 2018, i.e. telecom-resource suspension.

2. F2 (P0) — the market axis exists in the schema, the UI, and nowhere else

Three columns, zero readers

`orgs` carries `country_code`, `currency`, `timezone`. Grep across the whole API for their consumers:

- `packages/api/src/app/routes.ts:257-258`, `:369-370` — echoed in `GET /me`.
- `packages/api/src/app/routes.ts:311` — accepted in `PATCH /api/app/settings`.
- `packages/api/src/voice/org-queries.ts` — selected into a result object.
- `packages/web/src/web/pages/app/settings.tsx:64` — the UI defaults `countryCode` to `"IN"`.

Zero reads anywhere under `packages/api/src/voice/`. The only hits in that directory are the unrelated `listAvailableNumbers(orgId, countryCode, ...)` parameter in `twilio-provisioning.ts:198` and `admin-routes.ts:192`, which is a Twilio search filter, not the org's market.

`orgs.currency` has the same shape. The one place a currency reaches a prompt is `workflows/variables.ts`, and that reads `context.currency` from the *workflow payload* (Shopify's cart currency), not the org's. So an org that declares itself Indian does not get rupees in its agent's mouth; a Shopify webhook does.

What decides jurisdiction instead

`packages/weeber-compliance/src/calling-window.ts:30` → `isIndianNumber(e164)` → if `+91`, apply `packs/india.ts` (09:00–21:00 IST, single `Asia/Kolkata` timezone); otherwise apply `packs/us.ts` (area-code → state → state-specific window). The limitation is acknowledged in that file's own comment.

For calling windows specifically this is defensible — the callee's local time is what the window is about. For everything else it is wrong, and it is why F9 happens:

DECISION	SHOULD DEPEND ON	ACTUALLY DEPENDS ON
Calling window	callee locale	callee prefix — correct
Which compliance pack's consent/licensing rules apply	the <i>business's</i> regulator	callee prefix
Which templates an org may enable	the business's market	nothing — all 9 are public to all
Currency spoken in a quote	the business's market	Shopify webhook payload, or nothing
Default language / provider	the business's market	per-config <code>language</code> , NULL 14/17 times
Which pricing plan the org is on	the business's market	not implemented at all

Why this is a P0 and not a nit

`pricing-lock-2026-07-18.md` decision #3 is explicit: *"Geo-differentiated pricing — India and Global are separate plans, not a currency toggle,"* with two distinct pricing pages, two tier tables (₹2,499/₹12,999 vs \$79/\$499), and different compliance framing per market. That decision was locked three weeks ago and **there is no field in the product to hang it on**. Billing cannot select a plan family, the dashboard cannot show the right currency, and no code path can behave differently for an Indian customer than an American one.

Three of four orgs have `country_code` NULL. Even if something read the column, it would read nothing.

3. F3 (P1) — the Indic layer is fully built and has never run

The chain, verified end to end

- 1. `RECOMMENDED_LANGUAGES` (`agent-frame.ts:57-70`) — 12 entries: `en`, `hi`, `mr`, `ta`, `te`, `kn`, `ml`, `bn`, `gu`, `pa`, `hinglish`, `multi` . Comment at `:53` : *"India-first list since that's the primary market."*
- 2. `SARVAM_PREFERRED_LANGUAGES` (`:86`) + `prefersSarvam()` (`:88`) route the 10 Indic/Hinglish codes to Sarvam when no explicit provider is set — guarded by `SARVAM_API_KEY` (`stt/index.ts:40` , `tts/index.ts:37`). ADR-060.
- 3. `buildLanguageInstructionBlock()` (`agent.ts:293`) returns `""` when `!language || language === "en"` . Separate branches exist for `multi` (`:295`) and `hinglish` (`:311`); otherwise it emits *"Conduct this entire call in `label`"* (`:322`) plus an explicit instruction not to switch even if the caller does (`:303-317`). Emitted as prompt segment `"language"` (`agent.ts:363` , `:426`).
- 4. Provider normalisation: `toSarvamLanguageCode()` maps `hinglish` → `hi-IN` , else `{language}-IN` (`stt/sarvam.ts:116` , `tts/sarvam.ts:17`). `toDeepgramNova3Language()` (`stt/deepgram.ts:44`) **rejects `hinglish` and `ml` outright**, warns, and falls back to multi-language mode.
- 5. Audited canned Hindi/Hinglish greetings exist for 5 insurance templates (`voice/insurance-greetings.ts`); any other language returns `undefined` and falls through to an LLM-generated greeting.

That is a well-built, well-documented, well-tested subsystem. Now the production side:

```
org_agent_configs: 17 rows
  language:      NULL on 14; set on 3 (one 'en', one 'hi', one other)
  persona_prompt: NULL/empty on 17
calls 16–25:      stt_provider_used = deepgram    (10/10)
                  tts_provider_used = cartesia    (9/10; call 21 = elevenlabs, failover_count 2,
degraded)
                  llm_provider_used = gateway    (10/10)
Sarvam:          0 calls, all-time
```

So: language NULL → buildLanguageInstructionBlock returns "" → **no language instruction is in the prompt at all** → the model follows the (English) persona text → prefersSarvam(null) is false → Deepgram + Cartesia. Every branch of the Indic work is downstream of a field that is empty in 82% of production configs.

The unit-economics consequence

pricing-lock-2026-07-18.md §"Final quick sanity check" prices India Starter (₹2,499 / 357 min) with a margin range of **43% (Deepgram+Cartesia, ~₹4/min) to 83% (Sarvam, ~₹1.2/min)**, and describes Sarvam as *"the actual India-language default."*

It is not the actual default. It is the *conditional* default, conditional on a field that is unset. The observed production stack is the ₹4/min one. **43% is not the floor of that range — it is the current number.** The good case in the locked model has occurred zero times.

This also matters for the Sarvam Startup Program credits: they subsidise a provider we do not call.

4. F4 (P1) — our default STT is bottom-tier for the market we call primary

Source: arXiv 2604.19151v2, "Voice of India." Closed-source benchmark: 306,230 utterances / 536.1 hours / 36,691 speakers / 15 Indian languages / 139 regional clusters, **unscripted telephonic speech** (i.e. our actual channel, not read-aloud corpora). Metric is OIWER, multi-reference, tolerant of spelling variation. **Paper-cited, not measured by us.**

Word error rate, selected — lower is better; — = language not supported by that system:

	AS	BN	BHO	GU	HI	HNE	KA	MAI	ML	MR	OR	PA	TA	TE	UR
Sarvam Audio (Tier I)	12.7	6.1	20.9	12.8	5.0	17.6	16.3	24.8	18.9	9.4	14.0	11.2	14.2	18.2	7.0
Deepgram Nova-3 (Tier III)	—	28.9	45.8	—	13.0	42.4	53.7	—	—	43.7	—	—	67.8	43.1	—
ElevenLabs Scribe v2	—	—	—	—	7.7	—	—	—	23.0	—	—	—	20.4	—	25.5

Findings that matter to us:

- **Sarvam Audio is Tier I — lowest WER in 13 of 15 languages.**
- **Deepgram Nova-3 is Tier III, the paper's "inadequate systems" band** (ranks 9–16, WER > 35%), with coverage gaps in 7 of 15 languages. The paper separately cites Nova-3 Odia at **89.8**. Tamil at **67.8** means roughly two words in three are wrong.
- **No system is consistently under WER 20 across all 15 languages** — the paper's own stated practical-usability threshold. Even Sarvam exceeds it on Bhojpuri (20.9) and Maithili (24.8).

- Systematic 19–21% male-speaker penalty. District-level WER ranges ~4% (Nainital) to 44% (Mannarakkat). Out-of-region migrants degrade to 55–65%. Models that look strong on FLEURS degrade substantially here — read-corpus scores do not transfer to telephony.
- AssemblyAI Universal exceeds WER 100 on several languages (transcription failure). GPT-4o-mini Transcribe hits 295.9 on Gujarati. Gemini 3 Pro is competitive on Hindi (6.0) and best on dialectal Bhojpuri/Chhattisgarhi.

Three consequences:

1. The provider in our default path is the worst-rated of the credible options on the channel and the market we say is primary. The best-rated one is wired, credited, and idle (F3).
2. Audit 14's transcript-correctness findings were all on **English** calls. If an Indic pilot starts on the current default, add a 13–68% WER floor on top of those. No amount of prompt work survives an STT layer that mis-hears two-thirds of the words.
3. **Any marketing claim of "10+ Indian languages" is not defensible at usable quality.** Defensible today, on this evidence and with Sarvam actually selected: `hi`, `bn`, `pa`, `ur`, `mr`, `or`, plus `en`. Tamil/Telugu/Malayalam/Maithili/Bhojpuri are above the usability threshold on **every** system in the benchmark. Scope the claim to the languages, and say "with human escalation" for the rest.

Caveat I will not paper over: these are one paper's numbers on one closed benchmark, and Sarvam is both the top scorer and an Indian lab with an interest in that result. The direction is consistent with ADR-060's own reasoning and with Deepgram's published language list, but the *magnitudes* should be reproduced on our own recorded audio before they go in a deck. That requires keys (§8).

5. F5–F6 (P1) — the language field and the disclosure that does not follow it

F5 — a free-text input that switches providers

packages/web/src/web/pages/app/agents.tsx:361 :

```
<input id={`lang-${row.templateKey}`} value={form.language} onChange={(e) => set("language",
e.target.value)}
  placeholder="en, hi, mr..." list={`langs-${row.templateKey}`} className={fieldCls} />
```

An `<input>` with a `datalist`, next to two `<select>`s for STT and TTS provider. `language-support.md` defends this deliberately — "open free text... the curated list is a starting point, not a fence" — and as a *schema* argument that is right. As a *UI* argument it is wrong, because:

- `prefersSarvam()` (`agent-frame.ts:88`) is an exact lowercase membership test. `"hi"` → Sarvam. **"hi-IN"** → **not Sarvam**. **"Hindi"** → **not Sarvam**. **"hin"** → **not Sarvam**. All three silently take the English-first stack, and `buildLanguageInstructionBlock` cheerfully instructs the model to "Conduct this entire call in hi-IN".
- Nothing validates against `RECOMMENDED_LANGUAGES` on save.
- The field also keys the disclosure lookup (§F6) and the Sarvam `${language}-IN` code construction, so a typo degrades recognition quality, voice identity, and the legal line at once.

The fix is not to remove free text — it is a `<select>` of the 12 curated codes plus an explicit "other (advanced)" escape that warns that provider smart-defaults will not apply. Keep the escape hatch; stop making it the default interaction.

F6 — disclosure coverage is 3 of 12

`packages/weeber-compliance/src/consent.ts:58` , `DISCLOSURE_TEXT_BY_LANGUAGE` — `en` , `hi` , `hinglish` . That is **all**. `resolveDisclosure` (`:91`) normalises the tag (`hi-IN` → `hi`) and falls back to the English default for anything unmatched.

Correction to an earlier working note: an earlier note in this lane recorded that `hinglish` was missing a disclosure key. That was wrong — the romanised Hinglish line exists at `:61` , with a correct rationale in the comment (a Hinglish agent's TTS speaks Hindi via `hinglish` → `hi-IN` , so the disclosure must match the romanised surface form). The real gap is the other eight languages.

So a Marathi, Tamil, Telugu, Kannada, Malayalam, Bengali, Gujarati or Punjabi agent — every one of which routes to a Sarvam voice in that language — opens the call by speaking the **English** legal disclosure through that voice. It is arguably worse than no localisation: it is the one sentence on the call with legal weight, delivered in a language the callee may not have been addressed in.

Second, smaller point, in the same file's comment: *"Both non-English lines are drafts pending human review before they're spoken on a real call."* They are the live path. Nothing gates them.

`packages/weeber-compliance` is a **STOP-AND-ASK** package per `project-brief.md` . Flagged, not touched. The decision to make is binary and belongs to you: either extend coverage to the remaining eight, or narrow `RECOMMENDED_LANGUAGES` to what is compliance-complete. Shipping a language picker that outruns the disclosure table is the one option that should be off the table.

6. F7–F9 — customisation, attribution, and the US script inside India orgs

F7 — "custom agents" is a services motion priced as SaaS

`agentTemplates` `INSERT/UPDATE` live only in `packages/api/src/voice/admin-routes.ts:358` , `:412` , `:453` (the last flips a template to `visibility=private` + `ownerOrgId` , per ADR-086/091). There is **no merchant-facing create path**. A customer can toggle a template on, set language/providers/tone, and write a `persona_prompt` — nothing more.

Production:

- `agent_templates` : 9 rows, **all** `visibility=public` , **all** `owner_org_id NULL` , all active. Insurance ×6 (`appointment-setter` 11,197 chars, `feedback-nps` 10,382, `final-expense-qualifier` 19,480, `lead-followup` 10,337, `policy-renewal` 10,827, `post-sale-welcome` 12,425), Shopify ×3 (`cart-recovery` 8,386, `cod-confirmation` 7,933, `feedback` 6,579). **Zero private/bespoke templates exist**.
- `org_agent_configs.persona_prompt` : **NULL or empty on 17 of 17**. No customer has ever written a prompt.

The private-template machinery from ADR-086/091 is correct and has zero users. And that is a strategy question, not a bug: if agent authoring is something we do per account, Weeber is a managed service and should be sold and staffed as one. The locked ₹2,499 / \$79 self-serve tiers assume the opposite. Pick one; they cannot both be true.

F8 — calls are not attributable

`\d calls` — **no** `template_key` , **no** `agent_config_id` , **no** `language` . It does carry `stt_provider_used` , `tts_provider_used` , `llm_provider_used` , `health_status` , `provider_failover_count` , `estimated_cost_usd_cents` , `sentiment` , `intent` , `disposition` .

So we can answer "which vendor served this call" but not "which agent, which template, which language." Which means every question raised in this document — does Sarvam actually reduce WER on our audio, do Hindi calls convert differently, which template is worth the 19k characters — is unanswerable from production data, now and retroactively. Three additive columns, one migration, zero risk under the additive-only invariant. This is the cheapest finding in the document and it gates the value of every other one.

F9 (P2) — the US script is live inside India-declared orgs

`docs/agent-prompts/09-insurance-final-expense-qualifier-agent.md` states its scope explicitly: *"US, English-only... unlike the 04–08 insurance agents, which are bilingual EN/HI for the India+US launch,"* and that it exists because a US final-expense agency wanted its existing closer script.

`insurance-final-expense-qualifier` — 19,480 characters, US-only, English-only — is **enabled on 2 of 3 insurance orgs** (config 24 on `krishn`, config 36 on `presistentads` with explicit `deepgram / cartesia / gateway`). That is the direct explanation for the USD burial-cost figures observed in calls 24 and 25.

There is no mechanism that could prevent this: `agent_templates` and `org_agent_configs` have **no region / market / jurisdiction column** (verified via `\d`). Template market scope exists only as prose inside `default_persona_prompt` and in `docs/agent-prompts/*.md`, which no runtime reads. Combined with F2, the market axis is absent at both ends — the org does not declare one in a way code can read, and the template does not declare one at all.

7. What we are wrong about

Not defects — beliefs the evidence contradicts.

1. "India-first" describes the roadmap, not the product. Every production number is a US Twilio DID. Every production call ran an English-first stack. The most elaborate template in the catalogue is US-only by design and enabled on two thirds of insurance orgs. The India insurance path is hard-blocked by our own correct gate with no route to unblocking on Twilio. What is deployed is a **US outbound insurance dialer with an unexercised India-capable subsystem attached**. That may well be the right product — it is the one with legal delivery, real margins, and 25 calls of evidence — but the docs, the prompts, and the pricing all describe a different one.

2. Compliance-first and India-first currently contradict each other. The compliance moat is real: DNC with no bypass, consent records, opt-out events, calling windows, FTSA attempt caps, producer licensing, the audit-trail export. Almost all of it is **US law**, and all of it runs on Twilio. India compliance is a *telephony relationship* problem — 1600-series allocation, DLT registration, per-call DLT logging — and we have the gate but not the relationship. Being compliance-first in India means being an Exotel customer before it means writing another check.

3. Our Indic model supplier is now our competitor, and the India price floor moved. Sarvam has opened its own conversational voice-agent platform ("Samvaad") to public self-serve access (Inc42 exclusive; also reported by BFSI Elets and startupfeed). We plan to resell a margin on top of the application layer of the company whose model we depend on. Separately, **Bolna raised \$6.3M (₹57.3 Cr) seed led by General Catalyst, announced 21 January 2026** (YC and Blume existing), positioned verbatim on "India's scale, linguistic complexity, and cost sensitivity." Gnani.ai is launching five new agents. India all-in retail sits at roughly ₹1–15/min with a ₹3–5/min budget tier and at least one operator reporting ₹2/min on Sarvam (search-level, vendor/community sourced, not primary). Our locked India Starter is ₹2,499/357 min $\approx$ **₹7/min effective** — above the budget tier we would be compared against, and its margin only works at all if Sarvam is the default, which F3 shows it is not.

4. Global pricing, by contrast, has real headroom. US/global voice AI runs \$0.05–0.10/min entry, \$0.08–0.30/min typical, \$0.25–0.50 managed, up to \$2.00 premium; an independent teardown puts full pipeline COGS near \$0.029/min. Our $\sim$ \$0.06/min

all-in is comfortably inside that. **The same COGS number that prices us out of India's budget tier is a healthy margin in the US.** The market that economics points at is not the market the code comments name.

5. Prompt architecture will not survive the matrix it is heading into. One template = one prose blob (up to 19,480 characters), with market, jurisdiction, currency, licensing and language rules expressed as English sentences inside it. The dimensions are vertical × market × language, and only vertical is a real field (ADR-031's vertical-agnostic seam). That is precisely why F9 happens, and it gets worse with every market added. The composition layer already exists — `agent.ts:363-429` builds the prompt from labelled segments — so the fix is to move market/jurisdiction/currency out of the blob and into segments resolved from data, not to rewrite nine templates.

6. One thing we are right about, worth defending. One fixed spoken language per call, no mid-call switching (`language-support.md` , ADR-060) is the correct call, for the reasons stated: voice identity, latency, mid-stream stability. Code-switching belongs in *understanding* (Deepgram `multi` , Sarvam `codemix`), and that is where it is. Do not let a competitor's "dynamic language switching" bullet reopen this.

8. What I could not verify

Stated plainly so nothing here gets over-read:

- **No provider keys in this environment.** `/home/user/.env` holds `GITHUB_TOKEN` only. Nothing in §4 was measured on Weeber audio; every WER figure is paper-cited and every price is vendor- or community-published.
 - **Production env is unconfirmed:** `LLM_PROVIDER` , `AI_GATEWAY_MODEL` , `GROQ_MODEL` , and critically **whether `SARVAM_API_KEY` is set in Railway.** `docs/voice-quality/language-support.md` asserts "Weeber prod: live as of 2026-07-19"; that is a doc claim, not an observation. If the key is *not* set, F3 is worse than described — the smart default is dead code in production rather than unreachable code.
 - **The deployed commit cannot be fingerprinted.** `api.weeber.ai` returns 404 `DEPLOYMENT_NOT_FOUND` ; there is no health/version route. All production-behaviour statements are inferred from `main` at 8498e58 plus DB rows.
 - India price-ceiling figures are search-result level, not primary. The TRAI Direction and the arXiv paper *are* primary and were fetched directly.
-

9. Recommended sequence

Ordered by what unblocks what, not by size.

1. **Decide the primary market for the next 90 days in writing, and let the code state it.** My read: **US-first for the pilot** — it is what is built, legal, and evidenced — with India as a deliberate second market gated on a telephony relationship. Whatever you decide, F2 must be closed either way. → **ADR-095.**
2. **Make market a first-class field** (ADR-095): `orgs.market` (+ optional per-agent override) driving compliance pack selection, template eligibility, prompt currency, default language, provider default, and pricing plan family. Keep the callee-prefix check — demote it from source of truth to a **mismatch detector** that flags an org dialling outside its declared market.
3. **Add `template_key` , `agent_config_id` , `language` to `calls` .** Additive, one migration, and it is the precondition for ever answering §3, §4, §6 or §7 with data.
4. **If India is in scope at all: get one Exotel BYO number with DLT + 1600-series registration on one org.** Until that row exists with `number_series = '1600'` , the India insurance vertical is a slide, not a product, and no amount of code changes that.

5. **Set language explicitly on all 17 configs; make the field a select.** Then default India-market insurance agents to `hi / hinglish` (which routes them to Sarvam) and re-measure — on our own audio, before quoting anyone's WER table.
6. **Resolve the disclosure-coverage decision (STOP-AND-ASK):** extend to the remaining eight languages, or narrow the picker to `en / hi / hinglish`. Also decide whether "drafts pending human review" should be spoken on real calls.
7. **Resolve F7 explicitly: services or self-serve.** If self-serve, the merchant authoring path is the gap and the pricing is coherent. If services, the pricing page is describing a product we do not sell.

No code changed. One ADR proposed (095, for item 2). Items 1, 4, 6 and 7 are yours to decide, not mine.

ADR-095: The callee's country code is not our customer's market

- Status: Proposed
- Date: 2026-08-10
- Supersedes: none
- Amends: none
- Related: ADR-031 (vertical-agnostic seam), ADR-060 (Indic smart provider default, mid-call switching rejected), ADR-072 (a provider contract is what the server accepts), ADR-086/ADR-091 (template visibility), ADR-093 (a vertical switch left the old vertical's agents live)
- Source: `audit/2026-08-10-audit-15-the-market-is-a-column-nobody-reads.md` (F1, F2, F5, F9)

Context

Weeber has a first-class **vertical** axis. `orgs.vertical` drives which templates an org sees, which metrics the dashboard renders, which terminology it uses, and which compliance gates fire — and ADR-031 made adding a vertical a data change rather than a code change. That seam works. ADR-093 even taught it to clean up after itself when an org switches.

It has no equivalent for **market**. And market is now doing real work in every decision document while existing nowhere in code:

- `agent-frame.ts:53`, in a comment on `RECOMMENDED_LANGUAGES`: *"India-first list since that's the primary market."*
- `pricing-lock-2026-07-18.md` decision #3, locked: *"Geo-differentiated pricing — India and Global are separate plans, not a currency toggle"* — two tier tables (₹2,499/₹12,999 vs \$79/\$499), two pricing pages, different compliance framing per market.
- `docs/agent-prompts/09-insurance-final-expense-qualifier-agent.md`: *"US, English-only... unlike the 04-08 insurance agents, which are bilingual EN/HI for the India+US launch."*

Three market-scoped decisions, expressed in a code comment, a strategy doc, and a prompt-authoring doc respectively. No runtime reads any of them.

What exists, and what reads it

`orgs` carries `country_code`, `currency`, `timezone`. Every consumer:

SITE	WHAT IT DOES WITH THEM
<code>app/routes.ts:257-258</code> , <code>:369-370</code>	echoes them in <code>GET /me</code>
<code>app/routes.ts:311</code>	accepts them in <code>PATCH /api/app/settings</code>
<code>voice/org-queries.ts</code>	selects them into a result object
<code>web/pages/app/settings.tsx:64</code>	renders the form, defaulting <code>countryCode</code> to <code>"IN"</code>

Zero reads under `packages/api/src/voice/`. The only hits in that directory are `listAvailableNumbers(orgId, countryCode, ...)` (`twilio-provisioning.ts:198`, `admin-routes.ts:192`), which is a Twilio number *search filter*, unrelated to the org's market. `orgs.currency` never reaches a prompt either — `workflows/variables.ts` reads `context.currency` off the workflow payload (Shopify's cart currency), not the org's.

They are write-only columns. In production, 3 of 4 orgs have all three NULL.

What decides jurisdiction instead

packages/weeber-compliance/src/calling-window.ts:30 — isIndianNumber(e164) → packs/india.ts (09:00–21:00 IST, single Asia/Kolkata window) if +91 , packs/us.ts (area-code → state → per-state window) otherwise. The file's own comment acknowledges the limitation.

For calling windows this is correct and must not change — the window is a statement about the callee's local time, so the callee's number is the right input. The defect is that this single inference is the only market signal in the system, so it has been silently promoted into decisions it cannot make:

DECISION	CORRECT INPUT	CURRENT INPUT
calling window	callee locale	callee prefix — correct, keep
which regulator's consent/licensing rules bind us	the business's regulator	callee prefix
which templates an org may enable	the business's market	nothing; all 9 templates are public to all orgs
currency in a spoken quote	the business's market	Shopify webhook payload, or nothing
default language and STT/TTS provider	the business's market	per-config language , NULL on 14/17 rows
which pricing plan family applies	the business's market	not implemented

The two failures this already causes in production

Templates leak across markets. agent_templates and org_agent_configs have no region / market / jurisdiction column (verified via \d), so a template's market scope exists only as prose inside default_persona_prompt . Result: insurance-final-expense-qualifier — 19,480 characters, explicitly US-only and English-only — is enabled on 2 of 3 insurance orgs, including the one org that has actually declared country_code = IN / timezone = Asia/Kolkata . It is the direct explanation for the USD burial-cost figures in calls 24 and 25. No mechanism could have prevented it: the org cannot declare a market in a form code reads, and the template does not declare one at all.

A locked pricing decision has nothing to attach to. Geo-differentiated plans were decided three weeks ago. Billing cannot select a plan family, the dashboard cannot pick a currency, and no code path can behave differently for an Indian customer than an American one — because there is no field that says which one an org is.

The adjacent finding that this ADR does not fix

insurance-gates.ts:74 requires an active number_series = "1600" row for any insurance org dialling India (TRAI Direction 16 Dec 2025, F. No. G-6/(8)/2025-QoS-Part(I), clause (iii) — consent does not cure it; the 15 Feb 2026 deadline has passed). All four production numbers are US Twilio DID's with number_series NULL, and Twilio does not allocate 1600-series numbers. That gate is correct and stays. It is a telephony-relationship gap (Exotel/DLT), not a schema gap, and no decision in this ADR changes it. It is recorded here only because a declared market is what makes the gap visible before dial time instead of at it.

Decision

Market becomes a first-class, org-owned field that the runtime reads. The callee's country code is demoted from source of truth to cross-check.

- 1. orgs.market — a small closed enum ("IN" | "US" to start; additive-only, grows the way vertical does), set at onboarding next to vertical, with the existing country_code / currency / timezone columns kept and derived-or-set

alongside it. Not renamed, not dropped (additive-only invariant). `market` is the field code reads; the legacy three stay for display and are backfilled.

2. **Optional per-agent override** on `org_agent_configs`, same precedence discipline as providers and language: an explicit per-agent value wins over the org value, which wins over the platform default. A multi-market agency org is a real shape and this is where it is expressed.
3. **`agent_templates.market`** — nullable, where NULL means market-agnostic. Template eligibility is then `vertical` **and** `market`, resolved in `visibleTemplatesForOrg / loadVisibleTemplate` (ADR-091's single by-key read path), so a US-only template cannot be enabled by an India-market org. Existing rows: `09-final-expense` → `"US"`, the rest → NULL until authored otherwise.
4. **`market` drives, from one resolution point:** the compliance pack for business-side rules (licensing, consent basis, disclosure requirements — *not* the calling window), template eligibility, currency in the prompt's facts block, the default `language`, the resulting STT/TTS provider default via `prefersSarvam`, and the pricing plan family. Emitted into the prompt as its own labelled segment (`agent.ts:363-429` already composes by segment), never as prose inside a persona blob.
5. **Callee prefix stays where it belongs.** `isIndianNumber` keeps deciding the calling window, unchanged. It additionally becomes a **mismatch detector**: dialling outside the org's declared market is logged and surfaced (and, for regulated verticals, gate-able) rather than silently reinterpreted as a different jurisdiction.

Consequences

What gets better. The US-only template stops being enableable by an India-market org, which is the actual observed defect. Geo pricing acquires something to key off. An India-market org gets Hindi/Hinglish and therefore Sarvam by default rather than by remembering to type a language code into a free-text box (audit 15 F5). The 1600-series gap becomes knowable at onboarding instead of at dial time.

What this costs. Three additive columns and a resolution function threaded into template listing, prompt composition, and provider defaulting. Existing rows need a backfill decision — `org_68497dd7` is `IN`, the other three are undeclared and dial US numbers, so the honest default is `"US"` with an explicit onboarding prompt rather than an inferred value. Every org's *visible* template list changes once step 3 lands; that is the point, and per ADR-093's precedent the rule is disable, never delete.

What is deliberately not decided here.

- Whether the primary market for the next 90 days is US or India. That is a business decision (audit 15 §9 item 1); this ADR makes either one expressible, and is worth doing under either answer.
- Any change to `packages/weeber-compliance`. Extending disclosure coverage beyond `en / hi / hinglish` (audit 15 F6) is a **STOP-AND-ASK** item and is not bundled in.
- Mid-call language switching. Still rejected, still correct (ADR-060). Market selects the language *before* the call; it does not make it dynamic.
- Whether agent authoring is self-serve or a service (audit 15 F7). Unrelated axis, separate decision.