

Morning Update — Weeber/openvent Engineering

Covering: Wednesday, July 16 → Thursday, July 17, 2026

| Executive Summary

Two days of work moved the insurance vertical from "prompt content only" to "compliance-gated and launch-ready for both India and the US," while also shipping a Global Compliance Engine foundation (consent ledger, jurisdiction-aware calling rules), a consent-management UI, real privacy/terms pages, and a research-backed GTM/reliability plan. **12 commits, zero regressions at any point** — the backend test suite grew from 353 to 382 passing tests without a single failure introduced. Two genuine pre-launch blockers (an already-passed India regulatory deadline, and a US producer-licensing gap) were identified through research and closed before any real insurance call has been placed.

| Key Findings

- **Compliance is now the product's structural backbone, not an afterthought.** Six Global Compliance Engine items shipped (bypass-proofing, versioned/localized disclosure, jurisdiction-pack architecture, a real US mini-TCPA fix, a purpose-scoped consent ledger), then two more insurance-specific dial-time gates on top of that foundation.
- **Two real bugs were found and fixed along the way**, not just features shipped: a zero-authentication security hole on workflow admin routes, and a tool-registry drift bug (`flagGuardrailEvent` referenced in scripts but missing from two templates' tool lists) — the second one caught by a new automated test written specifically to prevent recurrence.
- **The insurance vertical is now the most deeply built of the three live verticals** — 8 full agent prompts (vs. generic placeholders for Shopify/Clinic upsell scripts), dual India+US regulatory citations, and the only vertical with dedicated compliance gates

beyond the generic DNC/calling-window checks.

- **A real, dated compliance risk was caught and closed:** the TRAI/IRDAI 1600-series calling-number mandate for India had a February 15, 2026 deadline that had already passed with zero platform support for it — now built and gated.

| Timeline

Yesterday — Wednesday, July 16, 2026


- | | | |
|-------|---------|---|
| 14:01 | 24f6cf2 | Hindi/Hinglish voice support, Phase 4 (final phase) — agents tab recommendation UI |
| 14:31 | c8047a7 | Docs synced with the completed Hindi/Hinglish work |
| 15:31 | b4d18cd | Workflow Canvas analytics overlay + fixed a zero-authentication security hole on admin routes |
| 17:08 | 5fdabdc | Workflow Canvas doc cleanup — resolved a real contradiction between two research docs |
| 17:32 | 86b6d9b | Test baseline correction (ADR-056) — "38 failures" cited across prior sessions was a false signal |
| 18:09 | f973a5e | Global Compliance Engine Tier 0 — all 6 items (see below) |
| 18:16 | f2e9cbe | 3 more insurance agent prompts landed (appointment setter, post-sale welcome, feedback/NPS) |
| 18:34 | e250574 | Marketing pages + Consent Ledger UI — real /privacy//terms pages, fixed a marketing overclaim |
| 18:46 | 47cc3fc | Insurance India+US regulatory iteration across all 5 prompts + a real bug fix |

Today — Thursday, July 17, 2026

02:13 5ded139 **Insurance pre-launch blockers built** — India 1600-series gate + US producer-licensing gate

02:19 c85ccec Insurance launch readiness resource doc — DLT registration steps, advisor data, NIPR info

02:33 a176ec7 Session recap + infra/product resource inventory + naming audit + GTM/reliability research



| What Actually Shipped, By Theme

1. Global Compliance Engine — Tier 0 (all 6 items)

The foundational compliance work everything else this session builds on: 1. Closed the `BYPASS_COMPLIANCE` footgun — a client-supplied flag could disable every legal gate; now hard-disabled in production, never honored from a request at all. 2. Disclosure text + version now stored per call (audit-trail requirement). 3. Disclosure localized per call language (English + a draft Hindi/Hinglish line, pending sign-off). 4. Calling-window logic refactored from one hardcoded if/else into a jurisdiction-pack architecture (India, US) — adding a future jurisdiction is now a new file, not more branching. 5. A real US mini-TCPA gap fixed — Florida/Oklahoma/Washington now correctly cap calling hours at 8pm, not the federal 9pm baseline. 6. A generic, purpose-scoped consent ledger built (service/transactional/marketing/underwriting/ feedback) — the foundation the insurance gates and future jurisdiction work builds on.

2. Insurance vertical — from prompts to a launch-ready, compliance-gated product

- 3 new agent prompts landed (appointment setter, post-sale welcome, feedback/NPS), bringing the vertical to 8 full scripts.
- All 5 scripts hardened with dual India (IRDAI) + US (NAIC, state producer licensing) regulatory citations, and a new "replacement" guardrail (discussing switching/cancelling a policy is its own regulated topic, not just general advice).
- Two genuine pre-launch blockers researched and closed:
- **India:** the TRAI/IRDAI 1600-series calling-number mandate (deadline already passed) — a dial-time gate now blocks an insurance org from calling India without a registered 1600-series number.
- **US:** producer state licensing — a dial-time gate now blocks a call/transfer to a US lead in a state no licensed advisor on file covers, resolved via a new ~230-entry area-code-to-state map.
- Both ship as manual-entry today; researched NIPR (the real industry-standard producer-license database every compliance vendor in this space wraps) as the upgrade path, schema already shaped to accept it without another migration.

3. Marketing + product UI

- Real `/privacy` and `/terms` pages — previously dead links in the site footer.
- A genuine marketing overclaim found and fixed: the landing page and FAQ both said "every number is checked against your consent records" — true for the always-on Do-Not-Call check, not yet true for the new purpose-scoped consent system. Wording corrected rather than left overclaiming.
- Admin and merchant UI built for the consent ledger (search, summary cards) and the two new insurance compliance gates (number-series picker, licensed-advisor manager).
- Insurance added to the public-facing vertical grid (draft copy, flagged for sign-off).

4. Hindi/Hinglish voice support — completed (all 4 phases)

Finished the multi-phase effort to make Hindi/Hinglish agents genuinely good rather than "technically has a language dropdown entry" — live-verified against real ElevenLabs and Sarvam accounts and real Hinglish audio, catching 2 real silent bugs along the way that would otherwise have shipped broken.

5. Research delivered

- A full India + US regulatory research pass for insurance voice AI (IRDAI, NAIC, TCPA, state producer licensing, recording-consent-state variance).
- A GTM + reliability report: competitor pricing benchmarks (\$0.05–\$0.33/min effective across Retell/Vapi/ElevenLabs), the reliability bar to match (Retell's cross-provider-failover claim), onboarding-speed benchmarks, and SaaS trial-to-paid conversion data — with 5 prioritized, concrete recommendations, picked up next.

| Verification (every claim below checked this session, not carried from memory)

- `bun run typecheck` — clean across all 3 packages, every commit.
- `bun run lint` (oxlint) — 0 warnings, 0 errors, every commit.
- `bun run test` — **353 → 382 passing, 0 failures at any point** across the session.
- `bun run build` — succeeds, every commit where the frontend changed.

| What's Next

Per the agreed order: move to the 5 prioritized recommendations from the GTM/reliability research (`docs/product-infra-and-gtm-report.md`, Part 4) —

1. Cross-provider STT/TTS/LLM failover (the top reliability differentiator).
2. A deliberate pilot-vs-trial GTM decision for the compliance-heavy verticals.
3. Use the real competitor pricing range as an input when founder pricing is finalized.
4. Keep onboarding-speed marketing claims honest against realistic benchmarks.
5. Clean up the small naming items found in the audit (`PUBLIC_MERCHANT_APP_URL`, etc).

| Methodology

Internal engineering activity report — sourced directly from this repository's git history (`git log`), CI verification output (`typecheck/lint/test/build` run fresh this session, not recalled), and the session's own delivered docs (`changelog.md`, `docs/global-compliance-engine-plan.md`, `docs/agent-prompts/00-insurance-regulatory-reference.md`, `docs/insurance-launch-readiness.md`, `docs/product-infra-and-gtm-report.md`). Not a market research report — no external citations needed for the timeline/verification sections; the GTM section's claims are sourced in `product-infra-and-gtm-report.md` directly.