

The gate is on two of the five doors

US-insurance pilot readiness — the compliance chokepoint, the empty tables behind it, and the market pivot that the codebase has not been told about. Audited against the code, the production database, and primary regulatory sources.

DATE	10 August 2026
REPO / COMMIT	Aurora-091/openvent · main @ 8498e58
DATA	Production Postgres (Supabase, pooled) — 4 orgs, 4 phone numbers, 11 calls all-time, 4 leads, 0 advisors, 0 DNC rows, 0 consent records
CONTENTS	Audit 16 (7 findings) · ADR-096 (proposed) · ADR-097 (proposed)
STATUS	Findings only — no code changed, nothing committed

Anchoring note. This document is written after the market pivot was confirmed: no India pilot was ever landed, and one real US-licensed insurance agency is in pilot dialling US prospects with roughly five more on a waitlist. It therefore supersedes the India/Shopify-first framing in docs/brain/project-brief.md rather than checking work against it. Legal figures are cited to the FCC ruling and the 11th Circuit opinion; the filing-volume figure is vendor-tracked, not a primary docket count, and is labelled as such. Production behaviour is inferred from main plus DB rows — the deployed commit cannot be fingerprinted (no health route).

Audit 16 — the gate is on two of the five doors

Audit 15 was written on the assumption in `docs/brain/project-brief.md`: ecommerce, Shopify first, India-first. That assumption is now dead. The stated position as of today is: **one real US-licensed insurance agency in pilot, dialing US prospects, plus roughly five more on a waitlist. No India pilot was landed.** This audit re-runs pilot readiness against *that* reality instead of the brief.

The conclusion is not "we are ready and the market changed." It is that the pivot moves Weeber from a jurisdiction where our exposure is regulatory (TRAI suspends telecom resources) into one where the exposure is **private right of action at \$500–\$1,500 per call with no aggregate cap** — and it does so while the code has three defects that are individually survivable in a demo and jointly fatal in a real campaign.

Everything below is read from source at `8498e58` or queried from the production database. Nothing is inferred from the brief.

0. The headline

#	FINDING	SEV
F1	<code>placeOutboundCall</code> has five call sites and compliance gates run at two of them. The three ungated paths include <code>POST /api/app/leads/:id/call-now</code> — the tenant-facing leads dialer, i.e. the exact endpoint a pilot uses to call a lead. Ungated means no DNC, no calling window, no FTSA attempt cap, and no insurance gates. <code>project-brief.md</code> 's invariant "DNC has no bypass anywhere, on purpose" is false in code .	P0
F2	<code>insurance_advisors</code> is empty in production . <code>checkInsuranceProducerLicensing</code> blocks any <code>+1</code> call whose area code resolves to a state with no licensed advisor on file — so on the <i>gated</i> path the pilot's first real US campaign blocks 100% of resolvable calls, while the <i>ungated</i> leads path dials every one of them. The compliant door is shut and the non-compliant door is open.	P0
F3	<code>do_not_call</code> = 0 rows . <code>consent_records</code> = 0 rows . For US outbound insurance lead calling, the consent record is the TCPA defense and DNC scrubbing is the duty of care. Both tables exist, are wired, and are empty.	P0
F4	Zero US calls have ever been placed. All 11 calls all-time dialed <code>+91</code> from US Twilio DIDs. US English on Deepgram, area-code → state resolution, four-timezone calling windows, and Twilio AMD against US carriers are all completely unexercised in production. The pivot's target market has never received a call from this system.	P1
F5	Producer licensing is decided from the area code , not the lead's state. US mobile numbers are portable and routinely retained across state moves, so area code is not a defensible basis for a licensing determination. ADR-087 already added <code>state</code> to the lead intake schema for exactly this class of problem, and this gate ignores it — the only lead with any enrichment carries <code>{"city": "Newyork"}</code> and no <code>state</code> at all.	P1
F6	Nothing can be billed. <code>plan_name</code> is NULL on 4 of 4 orgs, there is no plan enforcement (<code>pricing-lock-2026-07-18.md</code> caveat), and Stripe rejected the Indian entity. The pivot's whole thesis is that US willingness-to-pay is 3–5x India's, and there is no path from a US customer's card to our account.	P1
F7	Audit 15's F1 (India insurance legally undeliverable, TRAI 1600-series) is descoped, not fixed . The gate stays correct and stays unsatisfiable on Twilio. Any future India insurance conversation reopens it unchanged.	P2

1. F1 (P0) — five doors, two guards

`place-outbound-call.ts:86` states the contract explicitly:

```
* Compliance gates are the caller's responsibility
* (both call sites already run them before reaching here); this only places
* the call and returns the session key to store state under.
```

"Both call sites" was true when written. It is now wrong by three. Every caller of `placeOutboundCall`, from `rg`:

CALL SITE	ENDPOINT	GATED?
<code>voice/workflows/scheduler.ts:119</code>	scheduled-call sweep	yes — 6 gates at <code>:81 – :85</code>
<code>voice/routes.ts:313</code>	<code>POST /api/voice/calls/outbound</code>	yes — DNC + window + 2 insurance + India series at <code>:283 – :303</code>
<code>app/routes.ts:957</code>	<code>POST /api/app/leads/:id/call-now</code>	no
<code>app/routes.ts:645</code>	tenant agent preview test-call (<code>userOrgId</code> auth)	no
<code>voice/routes.ts:840</code>	<code>POST /orgs/:orgId/agent-configs/:templateKey/test-call-phone</code> (admin key)	no

`app/routes.ts:946-957` is the one that matters. Its own comment says it dials "through the same outbound path as the manual test-call, honoring the org's outbound routing" — which is accurate about *routing* and silent about the fact that the manual test-call route it is imitating runs six gates first and this one runs none:

```
.post("/leads/:id/call-now", async (c) => {
  const orgId = c.get("userOrgId");
  const id = Number(c.req.param("id"));
  if (!Number.isFinite(id)) return c.json({ error: "invalid id" }, 400);
  const result = await getOrgLead(orgId, id);
  if (!result) return c.json({ error: "not found" }, 404);
  if (!isValidE164(result.lead.phone)) {
    return c.json({ error: "Lead phone is not a valid E.164 number, can't place a call." }, 400);
  }
  const placed = await placeOutboundCall({ orgId, to: result.lead.phone });
```

E.164 shape is validated. Legality is not.

Why this is the same defect as ADR-090's eight

`voice/routes.ts:279` carries a long comment about how thoroughly the `bypassCompliance` request-body flag was removed — "it's stripped entirely, not just gated," "the env var is now hard-disabled in production." That work was real and it hardened **one** of five doors. The bypass that actually exists in production today is not a flag; it is a second endpoint that never had the gates in the first place. This is the reachability class of defect that ADR-090 stood up the `knip` ratchet for, and `knip` cannot see it — the gate functions *are* imported and *are* called, just not on every path.

Severity under the pivot

In India the ungated path risks TCCCPR action against telecom resources. In the US, an ungated call to a number on the federal DNC registry is \$500–\$1,500 statutory, per call, no cap, with a private right of action and no need to prove harm — and the FCC's February 2024 declaratory ruling puts AI-generated voices squarely inside "artificial or prerecorded voice," so every call our product makes is in scope. TCPA class filings are running well above last year's pace (a record 211 in February 2026 per TCPAWorld's tracker; vendor/tracker-cited, not a primary docket pull). A 10,000-lead campaign through `call-now` with no DNC scrub is a \$5M–\$15M theoretical exposure on one tenant's afternoon.

The multi-tenant shape makes it worse than a single-company mistake: the pilot is licensed and presumably careful, but the waitlist is five unknown agencies and the ungated endpoint is available to every one of them from day one.

This is the single thing to fix before the pilot runs a campaign. See ADR-096.

2. F2 (P0) — the compliant path blocks every US call

```
insurance-gates.ts:103 checkInsuranceProducerLicensing :
```

- no-op unless `to` is NANP-shaped and `orgs.vertical === "insurance"` ✓ (pilot is both)
- honors the self-expiring 24h `callingWindowTestModeUntil` bypass — **all four orgs' windows have expired** (`krisn`'s lapsed 2026-08-04; the other three are NULL)
- `resolveUsState(toNumber)` → if unresolved, **fails open** and allows
- otherwise requires some row in `insurance_advisors` for this org whose `licensedStates` contains that state

Production:

```
select org_id, count(*) from insurance_advisors group by 1; → (0 rows)
select id, org_id, name, licensed_states from insurance_advisors; → (0 rows)
```

`advisors.some(...)` over an empty array is `false`. So for every US lead whose area code resolves to a state — which is most of them — the gated path returns 403 with "No licensed advisor on file for this org is licensed in {STATE}". The pilot is a licensed agency; their licences exist in the real world and in no row of our database, and there is no evidence any onboarding step asks for them (`assigned_advisor_id` is NULL on all 4 leads).

The perverse outcome

Put F1 and F2 together and the product actively teaches the pilot to use the illegal path:

- `POST /api/voice/calls/outbound` → **403, blocked**, correctly, for a paperwork gap
- `POST /api/app/leads/:id/call-now` → **dials**, no DNC, no window, no licensing

A pilot who hits the 403, gets no clear remediation UI, and finds that the leads screen's "Call now" button just works will use the button. We will have built a compliance system whose only observable effect on a paying customer is to route them around itself.

3. F3 (P0) — the defense file is empty

```
select count(*) from do_not_call;      → 0
select count(*) from consent_records; → 0
```

`checkOutboundCallCompliance(to, dncAdapter)` runs on the gated path and is scrubbing against nothing. The federal DNC registry is not loaded; there is no internal suppression history because no US call has ever been made; and `consent_records` — the table that would hold the "prior express written consent" that makes an AI-voice insurance call lawful — has no rows, so today no call this platform places could be defended on consent-of-record.

Two things follow that are product decisions, not bugs:

1. **Where does consent come from?** The pilot buys or generates leads. Their consent lives in the lead vendor's or their own web form's records, not ours. If we do not ingest it at lead-import time and bind it to the phone number, our `consent_records` table is decorative and the pilot's TCPA defense sits in a spreadsheet somewhere else. ADR-089 already made CSV import previewable; consent provenance is the obvious next column, and it does not exist.
2. **DNC scrubbing is a subscription, not a feature.** Federal DNC registry access is per-area-code licensed and priced; internal + wireless + litigator lists are commercial feeds. Nobody has scoped or bought this. "Compliance-first is the moat" is currently an architecture, not a dataset.

4. F4 (P1) — the target market has never been dialed

All 11 production calls, all-time:

presistentads	outbound	+917905681369	completed	healthy	319s	2026-08-10 16:10
presistentads	outbound	+916393688162	completed	healthy	237s	2026-08-10 15:04
krisn	outbound	+919359848364	completed	healthy	89s	2026-08-10 10:42
krisn	outbound	+919359848364	completed	healthy	67s	2026-08-10 10:40
presistentads	outbound	+916393688162	completed	degraded	54s	2026-08-09 17:06
krisn	outbound	+919359848364	completed	healthy	19s	2026-08-09 13:41
rishipawar8999	outbound	+919359846364	completed	healthy	20s	2026-08-08 17:30
krisn	outbound	+919359848364	completed	healthy	20s	2026-08-08 16:24
krisn	outbound	+919359848364	completed	healthy	22s	2026-08-06 10:48
krisn	outbound	+919359848364	completed	healthy	19s	2026-08-06 10:46

`to_number` is `+91` on every row. `scheduled_calls` is empty — no campaign has ever run, only one-off dials. The two calls today at 5:19 and 3:57 are the only conversations of real length in the system's history.

This is fine as testing. It is not fine as evidence. Everything the pivot depends on is untested in production:

- **Deepgram Nova-3 on US English telephony** — likely our best case, but unmeasured here. Audit 15 had paper evidence it is Tier III on *Indian* speech; that finding is now largely irrelevant, and we have no equivalent evidence either way for the market we are entering.
- **resolveUsState on real lead data** — never executed against a resolvable US number in prod.
- **US calling windows across four timezones** — `packs/us.ts`'s area-code → state → timezone map has never gated a real call. India needed one timezone; the US needs the map to be right.
- **Twilio AMD** (`machineDetection: "DetectMessageEnd"`, `async`) — US answering-machine and carrier behavior is the environment it was written for and has never run there. Roughly half of US cold calls hit voicemail; if AMD misfires the agent talks to a machine for its full turn budget.

- **Latency as judged by a US buyer.** Audits 13 and 14 stand unfixed: the literal-greeting fast path is 0-for-11, p50 1539 ms of avoidable TTFT, and the agent does not know the current date. A US agency benchmarking against Vapi/Retell will hear that immediately, and the price umbrella that made India forgiving does not exist here.

Related but now demoted: all 11 calls dialed Indian consumers *from* US DID, which is the pattern TRAI/TCCCPR treats as unregistered telemarketing. As founder testing to known numbers it is a low concern; it is recorded here only so nobody later mistakes these rows for market evidence, and so the `presistentads` calls to two different `+91` numbers get an explicit "these were tests, right?" confirmation.

5. F5 (P1) — area code is not a state

`checkInsuranceProducerLicensing` resolves jurisdiction with `resolveUsState(toNumber)`. Its own doc comment concedes this is "best-effort." For a *calling window* that is defensible: guessing wrong shifts a call by an hour or two and `checkUsCallingWindow` already fails safe. For **producer licensing** it is the wrong test entirely — number portability means a `+1 212` mobile may have lived in Texas for a decade, and "we inferred the prospect's state from their area code" is not a story any carrier's compliance officer or state DOI will accept as the basis for who solicited whom.

The fix already half-exists. ADR-087 added `state` to the lead intake schema precisely because the opener needed it. The gate does not read it. And the intake is not populating it either — the only enriched lead in production is:

```
{ "city": "Newyork", "budget_band": "1k-3k", "existing_policy": "true",  
  "product_interest": "health", "best_callback_time": "Now", "preferred_language": "en" }
```

`city: "Newyork"` — free text, misspelled, no `state` key, and `leads.fields` is unvalidated `jsonb`. The correct precedence is stated lead state → area-code inference → block-or-flag, with the source of the determination recorded on the call so it is auditable afterwards. Note this makes licensing *stricter*, not looser: unresolved state currently fails open, and for a licensing question that is the wrong default once there is a real advisor roster to check against.

6. F6 (P1) — there is no way to take their money

`plan_name` is NULL on all four orgs. No plan/tier enforcement exists anywhere (an explicit caveat in `pricing-lock-2026-07-18.md`). Stripe rejected the Indian entity; Razorpay is weak on international card acceptance; Dodo/Paddle (merchant-of-record) were being evaluated and are not integrated.

The pivot's core financial claim is that Global Starter at \$79 / \$0.10 overage yields ~49–65% margin against India Starter's 43% floor. That claim is unrealisable until a US company can pay an Indian private limited for a SaaS subscription. `billing-metering-spec-2026-08-09.md` exists; the gateway decision does not. **This is now the gating commercial item, ahead of any feature**, because a pilot that goes well and cannot be converted is worse than no pilot.

7. What the pivot actually costs — for the record

The market decision itself is sound and is recorded separately in ADR-097. The honest ledger:

What US-first wins: 3–5x willingness-to-pay against the same dollar COGS; the only vertical where a customer has actually asked; a compliance-heavy buyer for whom our DNC/consent/licensing machinery is the product rather than overhead; a market where the 11th Circuit's January 2025 vacatur of the FCC one-to-one consent rule (*Insurance Marketing Coalition v. FCC*) left lead-gen consent *easier* than the FCC intended.

What it costs:

- `project-brief.md`'s "Launch vertical: ecommerce, Shopify first" is now false. So is the Shopify- → Woo/BigCommerce/Dukaan sequencing. Three of nine templates (cart recovery, COD confirmation, feedback) are parked, `weebbersh` is parked, and COD confirmation has no US market at all — COD is essentially absent from US ecommerce. That is not wasted work; it is work sequenced first that should have been sequenced second.
 - The entire Indic layer becomes dead capital in the near term. It already was — Sarvam ran on 0 of 11 calls — but the Sarvam Startup Program credits and the "10+ Indian languages" narrative now have no customer attached. Grants (IIMA/NSRCEL/SISFS/IIMK) are entity-scoped so eligibility survives, but we would be pitching an India-first Indic story to panels while billing only US customers, and a competent panel will ask. Decide now whether the grant narrative becomes "Indian company selling into the US" — which is a legitimate and well-understood story — rather than getting caught mid-pitch.
 - Liability shifts onto us. Platforms get named under direct-participation theories. Before any tenant dials 10,000 US numbers we need MSA indemnity, a consent-of-record requirement in the contract, and E&O cover. None of that is written.
 - Ops hours invert: US calling windows are roughly 6:30pm–7:30am IST. One-person on-call means every incident is a night incident.
 - We lose the price umbrella that made latency defects tolerable. Audits 13 and 14 stop being technical debt and become sales objections.
-

8. Sequence

Ordered by "what blocks the pilot's first real campaign," not by size.

1. **Gates to the chokepoint** (ADR-096). Move the six gates inside `placeOutboundCall` and make it fail closed, so no fifth door can be added later without them. Highest severity, smallest diff.
2. **Advisor roster onboarding + licensing precedence** (F2, F5). The pilot cannot make a legal US call until their licensed states are rows in our database, and the gate should read the lead's stated state before it guesses from an area code.
3. **DNC feed + consent provenance at import** (F3). Scope and buy the registry access; add consent source/timestamp/text to CSV import and refuse insurance-vertical imports without it.
4. **One real US call, then a real US campaign** (F4). Everything in §4 is unexercised. Do this before, not after, the pilot's first campaign.
5. **Billing gateway** (F6). Parallel track, no code dependency, but it gates revenue.
6. **Audits 13/14 fixes** — now customer-facing quality, not internal polish.
7. **ADR-095's `orgs.market`** — still the right structural fix, and cheaper now: with US-first decided, `market` has a real default and a real second value instead of being an abstraction.

Not in scope and deliberately parked: audit 15's F1 India 1600-series (unsatisfiable on Twilio, no India customer to satisfy it for), the Indic language work, and the Shopify/COD templates.

Appendix — production state as queried, 2026-08-10

```
orgs (4)
  krishna35672's workspace  shopify  country_code -  outbound -  plan -
  presistentads's workspace insurance country_code -  outbound +18573706834 plan -
  rishipawar8999's workspace insurance country_code IN  outbound +17754554413 plan -
  krisn                    insurance country_code -  outbound +17126257861 plan -

org_phone_numbers (4)  number_series NULL on all; 3 active + 1 released; all +1
calls (11)             all outbound, all to +91, 1 degraded, first 2026-07-18, last 2026-08-10
scheduled_calls        0 rows
leads                  4 rows (3 auto-created from calls, 1 manual with city="Newyork", no state)
insurance_advisors     0 rows
do_not_call            0 rows
consent_records        0 rows
callingWindowTestModeUntil expired on krisn (2026-08-04), NULL on the other three
```


ADR-096 — a gate on two of five doors is a door

Status

Proposed. No code changed. Raised by audit 16 F1/F2.

Context

`placeOutboundCall` (`packages/api/src/voice/place-outbound-call.ts:89`) is documented as the single call-placement entry point, and its doc comment assigns compliance to its callers:

Compliance gates are the caller's responsibility (both call sites already run them before reaching here); this only places the call and returns the session key to store state under.

That was true when written. It is now wrong by three. The callers are:

CALL SITE	ENDPOINT	GATED
<code>voice/workflows/scheduler.ts:119</code>	scheduled-call sweep	yes (6 gates)
<code>voice/routes.ts:313</code>	<code>POST /api/voice/calls/outbound</code>	yes (6 gates)
<code>app/routes.ts:957</code>	<code>POST /api/app/leads/:id/call-now</code>	no
<code>app/routes.ts:645</code>	tenant agent preview test-call	no
<code>voice/routes.ts:840</code>	admin <code>.../test-call-phone</code>	no

Ungated means no DNC, no calling window, no FTSA attempt cap, and none of the three insurance/India number-series gates. `docs/brain/project-brief.md` lists as a non-negotiable invariant that "DNC has no bypass anywhere, on purpose (ADR-003, ADR-007)". In code, three endpoints are a bypass.

Two facts make this urgent rather than tidy-up. First, the market decision recorded in ADR-097 moves the primary market to US insurance outbound, where an unscrubbed call carries \$500–\$1,500 statutory damages per call with no aggregate cap and a private right of action, and where the FCC's February 2024 declaratory ruling places AI-generated voices inside the TCPA's "artificial or prerecorded voice." Second, the one ungated *tenant-facing* endpoint is the leads dialer — the endpoint a pilot actually uses — and `insurance_advisors` being empty in production means the **gated** outbound route currently returns 403 for every resolvable US number while the **ungated** leads route dials it. The product as deployed routes a careful customer around its own compliance system.

This is the same defect class as the eight wiring defects that motivated ADR-090's `knip` ratchet, and `knip` structurally cannot catch it: the gate functions are imported and are called, just not on every path. Unit tests hide it for the same reason — each gate has tests, and the gates pass.

An additional live example of why caller-responsibility fails: `voice/routes.ts:279` carries a long comment about how completely the `bypassCompliance` request-body flag was removed ("stripped entirely, not just gated"). That hardening was real and it hardened one of five doors. The bypass that exists in production today is not a flag anyone shipped deliberately; it is an endpoint that was added later and never given the gates.

Decision

Compliance gates move inside `placeOutboundCall` and it fails closed. Placement stops being a primitive that trusts its callers and becomes the chokepoint that cannot be called around.

1. `placeOutboundCall` runs, before resolving routing or touching a provider: DNC, calling window, FTSA attempt cap, insurance number-series, insurance producer licensing, and the general India number-series flag — the same six, in the same order, that `voice/routes.ts:283-303` runs today.
2. It returns the existing typed failure shape (`{ ok: false, error, statusCode: 403 }`) so all five call sites surface a blocked call as a 403 with the gate's own reason string, unchanged in wording from today's messages.
3. The two currently-gated call sites **keep** their pre-checks. They are not removed. A duplicated gate is idempotent and cheap (both are DB reads on the hot path only for outbound placement), and the outbound route needs to reject before it writes a session. Removing them would trade one class of defect for another; the point of this ADR is that the floor is unconditional, not that the pre-checks were wrong.
4. **Fail closed on error.** If a gate throws — DB unavailable, adapter error — placement is refused, not allowed. This inverts today's implicit behavior in the ungated paths and is the whole point: the expensive failure is a call that should not have happened, not a call that did not happen.
5. The doc comment at `place-outbound-call.ts:80-88` is rewritten to say the opposite of what it says now, because that comment is what licensed the three ungated endpoints.
6. `BYPASS_COMPLIANCE` keeps its existing shape exactly — honored only outside production, never from a request body. Test mode (`orgs.callingWindowTestModeUntil`) keeps its existing scope exactly: calling window plus the two insurance config gates, self-expiring at 24h, and **never** DNC or the FTSA cap. This ADR does not widen or narrow any sanctioned bypass; it only makes the unsanctioned ones impossible.

The regression test that is the actual deliverable

A test that enumerates the callers of `placeOutboundCall` and asserts every one of them is covered is the only durable part of this decision — otherwise door six gets added in three weeks. Because `knip` cannot see this, the check is a source-level assertion in the api test suite: `placeOutboundCall`'s body must invoke the gate bundle before any provider client is constructed. Prefer one call to a single `assertOutboundCallAllowed(orgId, to)` helper so the assertion has one symbol to look for and the ordering lives in one place.

Consequences

Good. The invariant in `project-brief.md` becomes true. Adding a sixth dial path is safe by default. The pilot cannot be taught to route around compliance, because there is nowhere to route to. The 403 reasons become the single place remediation UI has to be written against.

Costs and risks.

- Both admin and tenant **test-call** endpoints become gated. This is intended, and it will immediately bite: with `insurance_advisors` empty, a test call to any resolvable US number will 403 until the pilot's licensed states are on file. That is the correct behavior and it makes audit 16 F2 a hard blocker rather than a latent one. The self-expiring test mode is the sanctioned escape for a live demo; if it is not enough, that is a separate, explicit decision about demo flows and not a reason to leave a hole.
- Six DB reads move onto the placement path for the three previously-ungated callers. Irrelevant for a scheduled or manual dial; noticeable only if someone loops `call-now`, which they should not.
- Duplicated checks on the two already-gated paths double those reads. Accepted for the reason in decision point 3.

- Fail-closed means a database blip stops all outbound dialing. Correct trade for this product.

What this does not decide. It does not change any gate's logic, does not touch `packages/weeber-compliance` (a STOP-AND-ASK boundary — the gate *functions* called here live in `packages/api/src/voice/compliance/`, and the compliance package is consumed unchanged), does not address the empty `insurance_advisors / do_not_call / consent_records` tables (audit 16 F2/F3), and does not fix the area-code-based licensing determination (audit 16 F5, which is a correctness question about *what* the gate reads, not *whether* it runs).

ADR-097 — the market that answered is not the market we planned

Status

Proposed. Requires the user's explicit acceptance, because it reverses a founding decision.

Context

`docs/brain/project-brief.md` states: *"Launch vertical: ecommerce, Shopify first (cart recovery + COD confirmation + feedback)"*, with clinic, insurance and hotel "on the board," and India as the primary market (`agent-frame.ts:53` , "India-first list since that's the primary market"; `pricing-lock-2026-07-18.md` decision #3 locking India and Global as separate plan families).

What actually happened over roughly three weeks of production existence:

- **No India pilot was landed.** Not ecommerce, not clinic.
- **One US-licensed insurance agency is in pilot**, dialing US prospects, with about five more agencies on a waitlist.
- Production reflects the insurance skew already: 3 of 4 orgs are `vertical=insurance` , all four outbound DID's are US Twilio numbers, and the 19,480-character explicitly-US-only final-expense template is `enabled` on 2 of 3 insurance orgs. The one shopify org has placed no calls.
- Nothing has been billed: `plan_name` NULL on 4 of 4.

Market evidence, independent of who happened to sign up:

- **US Shopify voice is the worst quadrant available**, so the original plan's US extension was never viable either. COD is effectively absent from US ecommerce, which removes the COD-confirmation agent entirely. Cart recovery competes against SMS at ~2.65% conversion for fractions of a cent, already installed on every US Shopify store via Klaviyo/Postscript — a \$0.10/min call must beat that by an order of magnitude. And an AI-voice cart-recovery call to a US mobile is a marketing call requiring prior express *written* consent under the FCC's February 2024 ruling; a checkout email field is not that consent.
- **India insurance is legally undeliverable as built.** TRAI's Direction of 16 December 2025 (F. No. G-6/(8)/2025-QoS-Part(I), primary PDF fetched) clause (iii) bars IRDAI-regulated entities from any service or transactional voice call from a non-1600-series number after 15 February 2026, and consent does not cure it. Twilio cannot allocate that series. Our own `checkInsuranceNumberSeriesCompliance` correctly blocks it and nothing on our stack can satisfy it.
- **India ecommerce is real but contested and thin.** COD/RTO pain is genuine and the 1600-series Direction does not cover retailers. But the effective ceiling is ~₹7/min against a locked 43% worst-case margin, and Bolna raised \$6.3M led by General Catalyst in January 2026 positioned explicitly on India's cost sensitivity, while Sarvam — our own Indic model supplier — has opened a self-serve conversational voice-agent platform, i.e. became a competitor at the application layer.
- **US insurance is the only quadrant with observed pull and margin headroom.** A customer asked. The locked Global Starter (\$79 / \$0.10 overage) yields ~49–65% margin on the same dollar COGS that leaves 43% in India. And the buyer is compliance-obsessed, which makes our DNC/consent/licensing/ disclosure machinery the product rather than overhead — the one place our existing investment is a differentiator instead of a tax.

Decision

US insurance outbound is the launch market and launch vertical. Everything else is parked, not deleted.

1. `project-brief.md` 's launch-vertical line is replaced: launch vertical is **insurance**, launch market is **US**. Ecommerce/Shopify, clinic, and hotel move to "on the board." The Shopify → Woo/BigCommerce/Dukaan sequencing is unscheduled.
2. **Nothing is deleted.** ADR-031's vertical-agnostic seam means the three Shopify templates, the Indic language layer, the Sarvam adapters, the India compliance pack, the 1600-series gates and `weeber.sh` all stay in the tree, tested, unreferenced by any active customer. Parking costs us almost nothing precisely because of ADR-031; deleting would cost us the option.
3. **India remains the second market, not an abandoned one** — reopened on either a real India ecommerce pilot or an India insurance customer who already holds 1600-series numbers. The preconditions are written down (below) so this is a decision to revisit, not a door slammed.
4. **ADR-095's `orgs.market` is now cheaper and should ship with this.** With US-first decided, `market` has a real default (`us`) and one real alternative (`in`) rather than being a speculative abstraction, and it is what lets the parked India work stay in the tree without leaking into a US customer's prompts — which is exactly how the US-only final-expense template ended up enabled on the `country_code=IN` org.
5. **The grant/investor narrative is explicitly restated** as "Indian company, Indian engineering cost base, selling into the US insurance market" — not India-first Indic voice. Grants (IIMA/NSRCEL/SISFS/IIMK, DPIIT) are entity-scoped so eligibility is unaffected, and the Sarvam Startup Program and NVIDIA Inception credits remain valid; what changes is that we stop telling a panel a customer story we do not have. Decide this deliberately now rather than being caught mid-pitch.

Preconditions before the pilot runs a real campaign

This ADR is a market decision, and it is worthless if the pilot's first campaign is unlawful or unbillable. In dependency order:

1. **ADR-096** — gates at the chokepoint. Blocking.
2. **Advisor roster onboarding** — `insurance_advisors` is empty, so producer licensing blocks every resolvable US number on the gated path. Blocking.
3. **DNC feed + consent provenance at lead import** — `do_not_call` and `consent_records` are both empty. Blocking for outbound at volume.
4. **One real US call, then one real US campaign.** No US number has ever been dialed by this system; US English STT, area-code → state resolution, four-timezone calling windows and Twilio AMD are all unexercised in production.
5. **A payment path that accepts a US company's card for an Indian private limited** (Stripe rejected; Dodo/Paddle undecided). Non-blocking for the pilot, blocking for revenue.
6. **MSA indemnity, contractual consent-of-record requirement, and E&O cover** before any tenant dials at volume. One bad tenant on the waitlist is \$5M–\$15M of theoretical statutory exposure and platforms get named under direct-participation theories.

Preconditions to reopen India

Any one of: an India ecommerce pilot with a signed paid commitment; an India insurance customer that already holds 1600-series numbers and DLT registration; or an Indian telephony partner (Exotel/ Ozonetel/similar) integrated to the point of placing a real call — `plivo-client.ts` and `exotel-client.ts` both still carry "no live prototype call yet" caveats.

Consequences

Good. Effort concentrates on the one quadrant with a customer, the best margin, and a buyer who values our compliance work. The latency and correctness defects in audits 13/14 get an owner, because a US agency benchmarking against Vapi/Retell will hear them.

Costs, stated plainly.

- Roughly the last month of Shopify/COD/Indic work has no customer attached. It was sequenced first and should have been second. The Indic layer was already dead capital in practice — Sarvam ran on 0 of 11 production calls — but this makes it explicit.
- We lose the price umbrella that made a 1.5s time-to-first-token survivable.
- Liability moves from regulatory (TRAI suspending telecom resources) to private and uncapped (\$500–\$1,500 per call, no aggregate cap, no need to prove harm). Our compliance architecture is good and our compliance *data* is empty; the gap between those two is now the main risk in the company.
- On-call inverts to Indian night hours.
- We have one pilot and a waitlist. A waitlist is not traction and should not be described as clients in any deck until money moves.

What this does not decide. It does not change `orgs.vertical` semantics or ADR-031, does not delete any template or package, does not set pricing (the locked doc stands, Global tiers simply become the active ones), and does not touch `packages/weeber-compliance`.